

VYSOKÉ UČENÍ TECHNICKÉ V BRNĚ

BRNO UNIVERSITY OF TECHNOLOGY

FAKULTA INFORMAČNÍCH TECHNOLOGIÍ
ÚSTAV POČÍTAČOVÉ GRAFIKY A MULTIMÉDIÍ

FACULTY OF INFORMATION TECHNOLOGY
DEPARTMENT OF COMPUTER GRAPHICS AND MULTIMEDIA

IMAGE CAPTIONING WITH RECURRENT NEURAL NETWORKS

SEMESTRÁLNÍ PROJEKT
TERM PROJECT

AUTOR PRÁCE
AUTHOR

Bc. JAKUB KVITA

BRNO 2015



VYSOKÉ UČENÍ TECHNICKÉ V BRNĚ
BRNO UNIVERSITY OF TECHNOLOGY



FAKULTA INFORMAČNÍCH TECHNOLOGIÍ
ÚSTAV POČÍTAČOVÉ GRAFIKY A MULTIMÉDIÍ

FACULTY OF INFORMATION TECHNOLOGY
DEPARTMENT OF COMPUTER GRAPHICS AND MULTIMEDIA

POPIS FOTOGRAFIÍ POMOCÍ REKURENTNÍCH NEU- RONOVÝCH SÍTÍ

IMAGE CAPTIONING WITH RECURRENT NEURAL NETWORKS

SEMESTRÁLNÍ PROJEKT

TERM PROJECT

AUTOR PRÁCE

AUTHOR

Bc. JAKUB KVITA

VEDOUcí PRÁCE

SUPERVISOR

Ing. MICHAL HRADIŠ, Ph.D.

BRNO 2015

Abstrakt

Výtah (abstrakt) práce v českém jazyce.

Abstract

Výtah (abstrakt) práce v anglickém jazyce.

Klíčová slova

Klíčová slova v českém jazyce.

Keywords

Klíčová slova v anglickém jazyce.

Citace

Jakub Kvita: Image Captioning with Recurrent Neural Networks, semestrální projekt, Brno, FIT VUT v Brně, 2015

Image Captioning with Recurrent Neural Networks

Prohlášení

Prohlašuji, že jsem tento semestrální projekt vypracoval samostatně pod vedením pana Michala Hradiše.

.....

Jakub Kvita
December 27, 2015

Poděkování

Zde je možné uvést poděkování vedoucímu práce a těm, kteří poskytli odbornou pomoc.

© Jakub Kvita, 2015.

Tato práce vznikla jako školní dílo na Vysokém učení technickém v Brně, Fakultě informačních technologií. Práce je chráněna autorským zákonem a její užití bez udělení oprávnění autorem je nezákonné, s výjimkou zákonem definovaných případů.

Contents

1	Introduction	2
2	Neural networks	3
2.1	Recurrent neural nets	3
2.1.1	LSTM – Long Short-Term Memory	4
2.1.2	GRU – Gated Recurrent Unit	5
2.1.3	Text sequences – Word level embeddings and character level	5
2.2	Convolutional neural nets	5
3	Experiments	6
3.1	Torch	6
3.1.1	nn, nngraph	6
3.1.2	rnn	6
3.1.3	Other packages	6
3.2	Predicting next character in sequence	6
4	Image caption generation	7
4.1	Related Work	7
4.1.1	Show and Tell	7
4.1.2	Show, Attend and Tell	7
4.1.3	From Captions to Visual Concepts and Back	7
4.1.4	Long-term Recurrent Convolutional Networks for Visual Recognition and Description	7
4.2	Datasets	8
4.2.1	MS COCO	8
4.2.2	Flickr 30k,8k	8
4.2.3	CIDEr datasets	8
4.3	Evaluation metrics	8
4.3.1	BLEU	8
4.3.2	CIDEr	8
4.3.3	METEOR	8
5	Model	9
5.1	Architecture	9
5.2	Training details	9
6	Conclusion	10

Chapter 1

Introduction

Klasicky popis toho co se tady bude dit, jak je to dulezite, atd.

Chapter 2

Neural networks

General idea of neural networks was slowly emerging after World War II. Perceptron, as a single neuron unit, was created in 1958 by Frank Rosenblatt¹, but became popular only after creation of backpropagation algorithm in 1975. At that time neural nets have not reached massive popularity, not because they are not working, but due to small computing power of machines back then and lack of datasets. Recently (after 2000) neural nets became popular again, rebranded as ‘Deep Learning’, because researchers realized that it is possible and very useful to stack neural nets on top of each other and create deep architectures, which are more practical than shallow ones. During this reinvention neural nets have been successfully applied in multiple fields like computer vision, speech recognition and natural language processing.

Since then various useful architectures and algorithms are now introduced almost every month. There is vast amount of various architectures and algorithms, in this chapter, I will describe only a couple – those used in this thesis.

2.1 Recurrent neural nets

Feedforward neural nets are extremely powerful models, which can be highly parallelized. Despite that, they can be only applied to problems with inputs and outputs, which have fixed dimensionality (e.g. one-hot encoding vectors). This is a serious drawback, as many of the real-world problems are defined as sequences with lengths that are unknown to us in beforehand. Soon recurrent neural networks were introduced and they proved to be very useful to this kind of task. There is vast amount of recurrent neural networks, many not suitable for sequential tasks like Hopfield network, which are very successful in specific tasks, but nevertheless not useful for us now.

Apart from classification, which can be more precise when using sequences, one of the most important tasks is next value prediction. This core task can be then extended very simply to predict arbitrary number of future values. Prediction problems are all around us, from the weather forecast and stock market prediction to the autocomplete in smartphones or web browsers.

We can understand recurrent neural networks as very deep forward nets with shared weights. It is called RNN unrolling and it is described in figure 1. Layers of this very deep net spread in time, together with the input sequence. This is very innovative idea,

¹The perceptron: A probabilistic model for information storage and organization in the brain. Rosenblatt, F. Psychological Review, Vol 65(6), Nov 1958, 386–408.




Figure 1: Unrolling of the recurrent neural net. (C. Olah 2015)

which enabled training RNN with backpropagation through time. It also shows that, as very deep networks, they have vanishing or exploding gradient problem, which means that the network is not able to learn long-term dependencies, even though in theory it should. This is a serious issue, which is caused by iterating many times over the weights and the activation function with derivatives > 1 (exploding gradient) or < 1 (vanishing gradient). Gradient then dies out and learning stops for distant dependencies. Among others this problem has been solved by the LSTM unit described in part 2.1.1, which is most popular now and following research resulting in GRU described in part 2.1.2.

2.1.1 LSTM – Long Short-Term Memory

Long Short-Term Memory nets are special kind of recurrent network, capable of learning long-term dependencies. This architecture was introduced by Hochreiter & Schmidhuber (1997) in [6] after prior research of vanishing gradient problem. Later architecture was refined and popularized by other researchers and nowadays LSTM is most used and popular RNN architecture used.

The LSTM unit designed that it can remember a value for an arbitrary length of time. It contains gates that determine when the input is significant enough to remember, when it should continue to remember or forget the value, and when it should output the value. All these gates are built from simple layers and mathematical operations, and have the same input - current network input, last output of the unit and hidden state value of the unit. Diagram of an LSTM unit is on the figure 2.




Figure 2: Unrolled LSTM unit. (C. Olah 2015)

Jak to vyresilo problem vyse. Pridat i rovnice, ktore pouzivam ja, rozebrat dopodrobna.

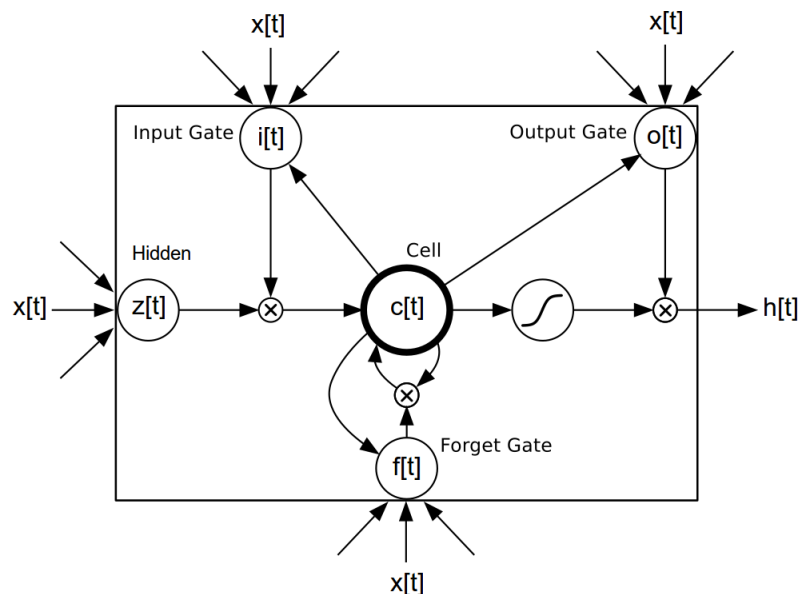


Figure 3: LSTM unit for equations

$$i_t = \sigma(W_x^i x_t + W_h^i h_{t-1} + W_c^i c_{t-1} + b^i) \quad (1)$$

$$f_t = \sigma(W_x^f x_t + W_h^f h_{t-1} + W_c^f c_{t-1} + b^f) \quad (2)$$

(1)
2

2.1.2 GRU – Gated Recurrent Unit

Zmíň jako updatovanou verzi

[2] [5]

2.1.3 Text sequences – Word level embeddings and character level

Možná trochu upravit název. (Character level and word level embeddings)

Popis toho jak se pracuje s textem v rnn, že to je taky sekvence. Character level, word level, embeddings. Popis rozdílu toho jak fungují překlady a generování dalšího prvku sekvence.

2.2 Convolutional neural nets

Kratký úvod do toho, kde se používají, jak se vyvinuly, jednoduchý popis toho jak fungují. Obrázek?

Asi není potřeba dávat subsekcce na vrstvy, stačí popsat jak to funguje všechno dohromady, jednotlivé vrstvy ve větách v jednom odstavci. Obrázek. V diplomě rozpracovat víc.

Chapter 3

Experiments

Kapitola jen na semestrální projekt. V diplomce ji odstraním.

Jak se to implementuje, jaké knihovny se používají - Caffe, Theano, TensorFlow, Torch. Popsat ze Torch bude v této kapitole.

Budu popisovat věci co jsem zkoušel implementovat v Torch.

3.1 Torch

Torch se zrecykluje do diplomky.

Udělat tedy tabulku o různých balících co torch má

Jak fungují rekurentní sítě v Torch.

Nacítání modelu z Caffe, ukládání v Torch...

[1]

3.1.1 nn, nngraph

Linky na knihovny v poznámkách pod čarou.

3.1.2 rnn

3.1.3 Other packages

loadcaffe, optim,...

3.2 Predicting next character in sequence

Jak jsem to udělal, co to dělá, ukázky.

Reference na Karpatyho char-rnn

[8]

Chapter 4

Image caption generation

Znovu uvod k tomu jak je to dulezite a tentokrat jak na tom lidi pracuji, co je potreba a jak se to hodnoti.

4.1 Related Work

Dat tomu nejake lepsi jmeno, clanky o popisovani obrazku ktere jsem cetl, pouzil.

4.1.1 Show and Tell

[14] [12]

Clanek z Coco od Googlu.

Zminit i strojovy preklad (Sequence to Sequence Learning with Neural Networks), architektura encoder, decoder

4.1.2 Show, Attend and Tell

[15]

Clanek z Coco z Montrealu/Toronta

4.1.3 From Captions to Visual Concepts and Back

[4]

Clanek z Coco od Microsoftu, mrknout se i na pokracovani v druhem clanku

4.1.4 Long-term Recurrent Convolutional Networks for Visual Recognition and Description

[3]

Clanek z Coco z berkeley

4.2 Datasets

COCO, Flickr, popis jake jsou. Asi zrusit sekce, udelat jen tabulku a mensi popis.

4.2.1 MS COCO

[10]

4.2.2 Flickr 30k,8k

[16] [7]

4.2.3 CIDEr datasets

[13]

4.3 Evaluation metrics

BLEU, cIDER, jak se používají, co dělají...

4.3.1 BLEU

[11]

4.3.2 CIDEr

[13]

4.3.3 METEOR

[9]

Chapter 5

Model

Do semestrálního projektu nebo až na diplomku?

Design modelu, co chci použít, jaké metody chci zkusit.

Položit si principiální otázku a zjistit, jestli to nějak pomůže, jak to funguje.

5.1 Architecture

Architektura modelu, jaké matematické modely jsem použil, bez implementačních detailů.

5.2 Training details

Popis pomocí jakého algoritmu jsme trénovali, s jakými parametry, minibatches, datasety.

Chapter 6

Conclusion

Udelat jeden zaver pro semestralni projekt, pak ho prepsat pro diplomku.

Bibliography

- [1] Torch – Scientific Computing for LuaJIT. <http://torch.ch/>. [Accessed: 2015-12-24].
- [2] Kyunghyun Cho, Bart van Merriënboer, Çağlar Gülçehre, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. Learning Phrase Representations using RNN Encoder-Decoder for Statistical Machine Translation. *CoRR*, abs/1406.1078, 2014.
- [3] Jeff Donahue, Lisa Anne Hendricks, Sergio Guadarrama, Marcus Rohrbach, Subhashini Venugopalan, Kate Saenko, and Trevor Darrell. Long-term Recurrent Convolutional Networks for Visual Recognition and Description. *CoRR*, abs/1411.4389, 2014.
- [4] Hao Fang, Saurabh Gupta, Forrest N. Iandola, Rupesh K. Srivastava, Li Deng, Piotr Dollár, Jianfeng Gao, Xiaodong He, Margaret Mitchell, John C. Platt, C. Lawrence Zitnick, and Geoffrey Zweig. From Captions to Visual Concepts and Back. *CoRR*, abs/1411.4952, 2014.
- [5] Klaus Greff, Rupesh Kumar Srivastava, Jan Koutník, Bas R. Steunebrink, and Jürgen Schmidhuber. LSTM: A Search Space Odyssey. *CoRR*, abs/1503.04069, 2015.
- [6] Sepp Hochreiter and Jürgen Schmidhuber. Long Short-Term Memory. *Neural Computation*, 9(8):1735–1780, November 1997.
- [7] Micah Hodosh, Peter Young, and Julia Hockenmaier. Framing Image Description as a Ranking Task: Data, Models and Evaluation Metrics. *Journal of Artificial Intelligence Research*, 47:853–899, 2013.
- [8] Andrej Karpathy. The Unreasonable Effectiveness of Recurrent Neural Networks. <http://karpathy.github.io/2015/05/21/rnn-effectiveness/>. [Accessed: 2015-12-30].
- [9] Alon Lavie and Abhaya Agarwal. METEOR: An Automatic Metric for MT Evaluation with High Levels of Correlation with Human Judgments. In *Proceedings of the Second Workshop on Statistical Machine Translation*, StatMT ’07, pages 228–231, Stroudsburg, PA, USA, 2007. Association for Computational Linguistics.
- [10] Microsoft. Common Objects in Context. <http://mscoco.org/dataset/#download>. [Accessed: 2015-12-29].
- [11] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. BLEU: A Method for Automatic Evaluation of Machine Translation. In *Proceedings of the 40th Annual Meeting on Association for Computational Linguistics*, ACL ’02, pages 311–318, Stroudsburg, PA, USA, 2002. Association for Computational Linguistics.

- [12] Ilya Sutskever, Oriol Vinyals, and Quoc V. Le. Sequence to Sequence Learning with Neural Networks. *CoRR*, abs/1409.3215, 2014.
- [13] Ramakrishna Vedantam, C. Lawrence Zitnick, and Devi Parikh. CIDEr: Consensus-Based Image Description Evaluation. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2015.
- [14] Oriol Vinyals, Alexander Toshev, Samy Bengio, and Dumitru Erhan. Show and Tell: A Neural Image Caption Generator. *CoRR*, abs/1411.4555, 2014.
- [15] Kelvin Xu, Jimmy Ba, Ryan Kiros, Kyunghyun Cho, Aaron C. Courville, Ruslan Salakhutdinov, Richard S. Zemel, and Yoshua Bengio. Show, Attend and Tell: Neural Image Caption Generation with Visual Attention. *CoRR*, abs/1502.03044, 2015.
- [16] Peter Young, Alice Lai, Micah Hodosh, and Julia Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the Association for Computational Linguistics*, 2:67–78, 2014.